

```
In [1]: 1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn import preprocessing,svm
6 from sklearn.model_selection import train_test_split
7 from sklearn.linear_model import LinearRegression
```

```
In [2]: 1 df=pd.read_csv(r"C:\Users\HP\OneDrive\Documents\fiat500_VehicleSelection_Dataset.csv")
2 df
```

Out[2]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

```
In [3]: 1 df=df[['engine_power','age_in_days']]
2 df.columns=['ep','aid']
```

```
In [4]: 1 df.describe()
```

Out[4]:

	ep	aid
count	1538.000000	1538.000000
mean	51.904421	1650.980494
std	3.988023	1289.522278
min	51.000000	366.000000
25%	51.000000	670.000000
50%	51.000000	1035.000000
75%	51.000000	2616.000000
max	77.000000	4658.000000

```
In [5]: 1 df.head(10)
```

Out[5]:

	ep	aid
0	51	882
1	51	1186
2	74	4658
3	51	2739
4	73	3074
5	74	3623
6	51	731
7	51	1521
8	73	4049
9	51	3653

In [6]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  --
 0    ep      1538 non-null    int64
 1    aid      1538 non-null    int64
dtypes: int64(2)
memory usage: 24.2 KB
```

In [7]: 1 df.fillna(method='ffill',inplace=True)

C:\Users\HP\AppData\Local\Temp\ipykernel_24420\4116506308.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.fillna(method='ffill',inplace=True)
```

In [8]: 1 x=np.array(df['ep']).reshape(-1,1)
2 y=np.array(df['aid']).reshape(-1,1)

In [9]: 1 df.dropna(inplace=True)


C:\Users\HP\AppData\Local\Temp\ipykernel_24420\1379821321.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.dropna(inplace=True)
```

In [10]: 1 *#Exploring the data scatter_plotting the data scatter*
2 sns.lmplot(x = "ep", y = "aid", data = df, order = 2, ci = None)


Out[10]: <seaborn.axisgrid.FacetGrid at 0x17f51b0ba00>



In [11]: 1 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
2 regr=LinearRegression()
3 regr.fit(x_train,y_train)
4 print(regr.score(x_test,y_test))


0.10327866026298738

```
In [12]: 1 y_pred=regr.predict(x_test)
2 plt.scatter(x_test,y_test,color='b')
3 plt.plot(x_test,y_pred,color='k')
4 plt.show()
```



```
In [13]: 1 df100=df[:][:100]
2 sns.lmplot(x='ep',y='aid',data=df100,order=1,ci=None)
3
```


Out[13]: <seaborn.axisgrid.FacetGrid at 0x17f378ce080>



```
In [14]: 1 df100.fillna(method='ffill',inplace=True)
2 X=np.array(df100['ep']).reshape(-1,1)
3 y=np.array(df100['aid']).reshape(-1,1)
4 df100.dropna(inplace=True)
5 X_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
6 regr=LinearRegression()
7 regr.fit(X_train,y_train)
8 print(regr.score(x_test,y_test))
9 print("Regression: ",regr.score(x_test,y_test))
10 y_pred=regr.predict(x_test)
11 plt.scatter(x_test,y_test,color='g')
12 plt.plot(x_test,y_pred,color='k')
13 plt.show()
```

0.3019208979888496

Regression: 0.3019208979888496



```
In [15]: 1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import r2_score
3 model=LinearRegression()
4 model.fit(X_train,y_train)
5 y_pred=model.predict(x_test)
6 r2=r2_score(y_test,y_pred)
7 print("R2_score: ",r2)
```

R2_score: 0.3019208979888496

Conclusion:

Dataset we have taken is poor for linear model but with the smaller data works well with linear model

```
In [ ]: 1
```